

Computational Linguistics for Indian Languages

(CS689A)

Assignment 2

SUBMITTED BY:

PRATIBHA GUPTA

231110038

QUESTION 2:

COMPARING MACRO F1-SCORE FOR INDIC BERT AND INDIC NER-

VALIDATION SET:

- MACRO F-SCORE FOR INDIC BERT: 0.6662
- MACRO F1-SCORE FOR 25 SENTENCES USING IndicNER: 0.7761

TEST SET:

- MACRO F-SCORE FOR INDIC BERT: 0.5357142857142857
- MACRO F1-SCORE FOR 25 SENTENCES USING IndicNER: 0.5614035087719298

From this we can observe that **IndicNER gives a better macro F1 score as compared to IndicBERT**. The higher F1-scores of IndicNER suggest that it's better suited for named entity recognition tasks, which aligns with its design and purpose. IndicNER have specialized architectures or training methodologies optimized for NER tasks, enabling it to capture nuances in named entity recognition better than IndicBERT. Since, IndicBERT is a strong language model; it might not be as effective in capturing specific NER-related features compared to a model explicitly designed for NER tasks like IndicNER. This observation suggests that although pre-trained language models like BERT can perform reasonably well across various NLP tasks, task-specific models like IndicNER might still outperform them in their respective domains.

QUESTION 4:

METRICS OUTPUT OF MANUALLY MARKED SENTENCES AND CHATGPT:

```
MACRO F1-SCORE FOR 25 SENTENCES:
0.47542807823312727
*****
CLASSWISE PRECISION, RECALL AND F1-SCORE FOR 25 SENTENCES:

O : [0.8168168168168168, 0.9477351916376306, 0.8774193548387097]
B_PER : [1.0, 0.7333333333333333, 0.846153846153846]
I_PER : [1.0, 0.8181818181818182, 0.9]
B_LOC : [0.7692307692307693, 1.0, 0.8695652173913044]
I_LOC : [0.3333333333333333, 1.0, 0.5]
B_ORG : [1.0, 0.16666666666666666, 0.2857142857142857]
I_ORG : [0, 0.0, 0]
B_MISC : [0.0, 0.0, 0]
I_MISC : [0.0, 0.0, 0]
```

METRICS OUTPUT OF MANUALLY MARKED SENTENCES AND TAGGING GIVEN BY Indic-BERT:

```
... MACRO F1-SCORE FOR 25 SENTENCES USING IndicBERT:
0.47190975693506293
CLASSWISE PRECISION, RECALL AND F1-SCORE FOR 25 SENTENCES:

O : [0.796969696969697, 0.9163763066202091, 0.852512155915721]
B_PER : [0.7368421052631579, 0.9333333333333333, 0.8235294117647058]
I_PER : [0.4166666666666667, 0.45454545454545453, 0.43478260869565216]
B_LOC : [0.25, 0.3, 0.2727272727272727]
I_LOC : [0.5, 1.0, 0.6666666666666666]
B_ORG : [0.4, 0.3333333333333333, 0.3636363636363636]
I_ORG : [1.0, 0.7142857142857143, 0.8333333333333333]
B_MISC : [0, 0.0, 0]
I_MISC : [0, 0.0, 0]
```

METRICS OUTPUT OF MANUALLY MARKED SENTENCES AND TAGGING GIVEN BY Indic-NER:

```
... MACRO F1-SCORE FOR 25 SENTENCES USING IndicNER:
0.2833128469036034
CLASSWISE PRECISION, RECALL AND F1-SCORE FOR 25 SENTENCES:

O : [0.7648809523809523, 0.8954703832752613, 0.8250401284109148]
B_PER : [0.3684210526315789, 0.4666666666666667, 0.4117647058823529]
I_PER : [0.16666666666666666, 0.18181818181818182, 0.17391304347826086]
B_LOC : [0.2222222222222222, 0.2, 0.2105263157894737]
I_LOC : [0, 0.0, 0]
B_ORG : [0.375, 0.5, 0.42857142857142855]
I_ORG : [0.6, 0.42857142857142855, 0.5]
B_MISC : [0, 0.0, 0]
I_MISC : [0, 0.0, 0]
```

QUESTION 5:

From the comparisons I can infer that accuracy of model and precision of predicting the tags highly depends on parameters like learning rate and batch size.

- Increasing the batch size decreases the precision of predictions made by the models.
- Also smaller learning rate gives better accuracy as compared with increasing the learning rate.

Hyper parameters setting to improve the performance of models:

The hyper parameters that I have changed:

- **per_device_train_batch_size and per_device_eval_batch_size** - These parameters define the number of training samples and evaluation samples, respectively processed simultaneously on each device (GPU or CPU) during training. Larger batch sizes require more memory, and if the batch size exceeds the available memory, training may fail. Smaller batch sizes generalize better and help prevent overfitting, especially with limited data.
- **Num_train_epochs** - It directly controls how many times the model iterates over the entire training dataset. Increasing this parameter allows the model to see the training data more times, potentially leading to better convergence and improved performance.
- **Learning_rate** - The learning rate determines the size of the step taken in the direction opposite to the gradient during optimization. A higher learning rate means larger steps, potentially leading to faster convergence but with the risk of overshooting the optimal solution.

Optimal values chosen by me:

Optimal results are obtained on following arguments: (for both the models)

- per_device_train_batch_size - 8
- Per_device_eval_batch_size - 8
- Num_train_epochs - 3
- Learning_rate - 5e-5

OUTPUT FOR BOTH THE MODELS -

OUTPUTS FOR INDIC-BERT (ARGUMENTS 1):

```
batch_size=8
args=TrainingArguments(
  output_dir='output_dir',
  per_device_train_batch_size=batch_size,
  per_device_eval_batch_size=batch_size,
  num_train_epochs=3,
  evaluation_strategy = "epoch",
  learning_rate=5e-5
)
```

```
***** eval metrics *****
epoch                =          3.0
eval_LOF_f1          =         0.7224
eval_LOF_number      =        10213
eval_LOF_precision    =         0.7169
eval_LOF_recall       =         0.728
eval_ORG_f1          =         0.5604
eval_ORG_number      =         9786
eval_ORG_precision    =         0.5707
eval_ORG_recall       =         0.5504
eval_PER_f1          =         0.7082
eval_PER_number      =        10568
eval_PER_precision    =         0.715
eval_PER_recall       =         0.7016
eval_loss            =         0.2606
eval_overall_accuracy =         0.9206
eval_overall_f1       =         0.6662
eval_overall_precision =         0.6705
eval_overall_recall   =         0.662
eval_runtime          =        0:04:19.63
eval_samples_per_second =        51.842
eval_steps_per_second =         3.243
```

OUTPUTS FOR INDIC-BERT (ARGUMENTS 2):

```
batch_size=6
args=TrainingArguments(
  output_dir='output_dir',
  per_device_train_batch_size=batch_size,
  per_device_eval_batch_size=batch_size,
  num_train_epochs=3,
  evaluation_strategy = "epoch",
  learning_rate=4e-5
)
```

```
***** eval metrics *****
epoch                =          3.0
eval_LOF_f1          =         0.7227
eval_LOF_number      =        10213
eval_LOF_precision    =         0.7183
eval_LOF_recall       =         0.7272
eval_ORG_f1          =         0.5614
eval_ORG_number      =         9786
eval_ORG_precision    =         0.5665
eval_ORG_recall       =         0.5564
eval_PER_f1          =         0.7064
eval_PER_number      =        10568
eval_PER_precision    =         0.7162
eval_PER_recall       =         0.6969
eval_loss            =         0.2626
eval_overall_accuracy =         0.9202
eval_overall_f1       =         0.6657
eval_overall_precision =         0.6693
eval_overall_recall   =         0.6621
eval_runtime          =        0:04:33.20
eval_samples_per_second =        49.267
eval_steps_per_second =         4.107
```

OUTPUTS FOR INDIC-BERT (ARGUMENTS 3):

```
batch_size=16
args=TrainingArguments(
  output_dir='output_dir',
  per_device_train_batch_size=batch_size,
  per_device_eval_batch_size=batch_size,
  num_train_epochs=3,
  evaluation_strategy = "epoch",
  learning_rate=5e-2
)
```

```
***** eval metrics *****
epoch                =          3.0
eval_LOC_f1          =          0.0
eval_LOC_number      =        10213
eval_LOC_precision   =          0.0
eval_LOC_recall      =          0.0
eval_ORG_f1          =          0.0
eval_ORG_number      =         9786
eval_ORG_precision   =          0.0
eval_ORG_recall      =          0.0
eval_PER_f1          =          0.0
eval_PER_number      =        10568
eval_PER_precision   =          0.0
eval_PER_recall      =          0.0
eval_loss            =         0.7806
eval_overall_accuracy =         0.8204
eval_overall_f1       =          0.0
eval_overall_precision =          0.0
eval_overall_recall   =          0.0
eval_runtime         =      0:03:59.08
eval_samples_per_second =        56.299
eval_steps_per_second =         1.761
```

OUTPUTS FOR INDIC-NER (ARGUMENTS 1):

```
args=TrainingArguments(
  output_dir='output_dir',
  per_device_train_batch_size=8,
  per_device_eval_batch_size=8,
  num_train_epochs=3,
  evaluation_strategy = "epoch",
  learning_rate=5e-5
)
```

```
***** eval metrics *****
epoch                =          3.0
eval_LOC_f1          =         0.8315
eval_LOC_number      =        10213
eval_LOC_precision   =         0.8116
eval_LOC_recall      =         0.8523
eval_ORG_f1          =         0.6783
eval_ORG_number      =         9786
eval_ORG_precision   =         0.6715
eval_ORG_recall      =         0.6853
eval_PER_f1          =         0.8122
eval_PER_number      =        10568
eval_PER_precision   =         0.8006
eval_PER_recall      =         0.8242
eval_loss            =         0.2078
eval_overall_accuracy =         0.9459
eval_overall_f1       =         0.7761
eval_overall_precision =         0.7635
eval_overall_recall   =         0.7891
eval_runtime         =      0:05:11.69
eval_samples_per_second =        43.183
eval_steps_per_second =         2.701
```

OUTPUTS FOR INDIC-NER (ARGUMENTS 2):

```
args=TrainingArguments(  
    output_dir='output_dir',  
    per_device_train_batch_size=8,  
    per_device_eval_batch_size=8,  
    num_train_epochs=3,  
    evaluation_strategy = "epoch",  
    learning_rate=5e-3)
```

```
***** eval metrics *****  
epoch = 3.0  
eval_LOF_f1 = 0.8315  
eval_LOF_number = 10213  
eval_LOF_precision = 0.8116  
eval_LOF_recall = 0.8523  
eval_ORG_f1 = 0.6783  
eval_ORG_number = 9786  
eval_ORG_precision = 0.6715  
eval_ORG_recall = 0.6853  
eval_PER_f1 = 0.8122  
eval_PER_number = 10568  
eval_PER_precision = 0.8006  
eval_PER_recall = 0.8242  
eval_loss = 0.2078  
eval_overall_accuracy = 0.9459  
eval_overall_f1 = 0.7761  
eval_overall_precision = 0.7635  
eval_overall_recall = 0.7891  
eval_runtime = 0:05:11.69  
eval_samples_per_second = 43.183  
eval_steps_per_second = 2.701
```

OUTPUTS FOR INDIC-NER (ARGUMENTS 3):

```
batch_size=8  
args=TrainingArguments(  
    output_dir='output_dir',  
    per_device_train_batch_size=batch_size,  
    per_device_eval_batch_size=batch_size,  
    num_train_epochs=3,  
    evaluation_strategy = "epoch",  
    weight_decay=0.1,  
    learning_rate=4e-5)
```

```
***** eval metrics *****  
epoch = 3.0  
eval_LOF_f1 = 0.8328  
eval_LOF_number = 10213  
eval_LOF_precision = 0.8136  
eval_LOF_recall = 0.853  
eval_ORG_f1 = 0.683  
eval_ORG_number = 9786  
eval_ORG_precision = 0.6776  
eval_ORG_recall = 0.6884  
eval_PER_f1 = 0.814  
eval_PER_number = 10568  
eval_PER_precision = 0.8022  
eval_PER_recall = 0.8262  
eval_loss = 0.2032  
eval_overall_accuracy = 0.9464  
eval_overall_f1 = 0.7787  
eval_overall_precision = 0.7668  
eval_overall_recall = 0.791  
eval_runtime = 0:05:07.58  
eval_samples_per_second = 43.76  
eval_steps_per_second = 2.737
```

OUTPUTS FOR INDIC-NER (ARGUMENTS 4):

```
batch_size=6
args=TrainingArguments(
  output_dir='output_dir',
  per_device_train_batch_size=batch_size,
  per_device_eval_batch_size=batch_size,
  num_train_epochs=3,
  evaluation_strategy = "epoch",
  learning_rate=5e-7)
```

```
***** eval metrics *****
epoch                =      3.0
eval_LOC_f1          =    0.5742
eval_LOC_number      =    10213
eval_LOC_precision    =    0.5671
eval_LOC_recall       =    0.5815
eval_ORG_f1          =    0.3439
eval_ORG_number      =     9786
eval_ORG_precision    =    0.3465
eval_ORG_recall       =    0.3413
eval_PER_f1          =    0.5299
eval_PER_number      =    10568
eval_PER_precision    =    0.4378
eval_PER_recall       =    0.6712
eval_loss             =    0.3268
eval_overall_accuracy =    0.9002
eval_overall_f1       =    0.4896
eval_overall_precision =    0.4509
eval_overall_recall    =    0.5356
eval_runtime          = 0:05:56.19
eval_samples_per_second = 37.788
eval_steps_per_second  =     3.15
```

CONCLUSION:

From the above outputs we can conclude that increasing the batch size significantly decreases the accuracy as smaller batch size gives better generalization and also helps reduce overfitting. And lower learning rate gives better convergence and helps models gain better accuracy. So, we can alter the hyper parameter for the training arguments and tune it to the most optimal one.